

# Markov Decision Process

Sunmook Choi `felixchoi@korea.ac.kr`

September 4, 2023

## 1 Markov Decision Processes

**Definition 1** (Markov Decision Process, MDP). *MDP is a tuple  $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$  where*

- $\mathcal{S}$  is the state space,
- $\mathcal{A}$  is the action space,
- $P$  is the transition probability,
  - \* the transition probability from the state  $s \in \mathcal{S}$  to the state  $s' \in \mathcal{S}$  given an action  $a \in \mathcal{A}$ .
  - \*  $P_{ss'}^a := P(S_{t+1} = s' | S_t = s, A_t = a)$
  - \*  $P(S_{t+1} = s' | S_t = s) = P(S_{t+1} = s' | S_t = s, S_{t-1} = s_{t-1}, \dots, S_1 = s_1, S_0 = s_0)$
  - \* We assume stationary Markov Process, that is,  $P_{ss'}^a$  does not depend on time step  $t$ .
- $R$  is the reward function, and
- $\gamma \in [0, 1]$  is the discount factor.

**Example 1** (Grid World). Let's consider an example, the Grid World, illustrated in Fig. 1.1. We call a robot agent which is moving in the grid world. The thing it interacts with, comprising




Figure 1.1: An example of Grid world Environment

everything outside the agent, is called environment. There are total 11 places that an agent can stay

in,  $(1, 1), (2, 1), \dots, (4, 3)$ , except for  $(2, 2)$ . Such places are called *states* of the grid world. The agent starts the game from the state  $(1, 1)$ , and the game ends when it reaches one of the final states,  $(4, 2)$  or  $(4, 3)$ . The agent can move either North, South, East, or West. The agent should take an *action* in order to move; the actions are North, South, East, and West. However, the grid world is *slippery* so that the next state is determined stochastically. The transition occurs corresponding to the transition probability  $P_{ss'}^a$ , which has the Markov Property, and an example is as follows:

$$P(s_{t+1} = (3, 2) | s_t = (3, 1), a_t = \text{North}) = 0.8$$

$$P(s_{t+1} = (2, 1) | s_t = (3, 1), a_t = \text{North}) = 0.1$$

$$P(s_{t+1} = (4, 1) | s_t = (3, 1), a_t = \text{North}) = 0.1.$$

Here, if the state transition is impossible due to the wall of the grid world, the agent stays in the current state. For each step, the agent gets reward from the environment. The agent gets small negative reward  $c$  (e.g.,  $c = -0.1$ ) at each step until it reaches the final state. This encourages the agent to reach the final state as soon as possible. The agent gets  $+1$  reward when it reaches the state  $(4, 3)$ , and gets  $-1$  reward when it reaches  $(4, 2)$ . The description above defines the state space, the action space, the transition probability, and the reward function. With a hyperparameter  $\gamma$ , the grid world is a Markov Decision Process.

The general interaction between an agent and an environment can be illustrated as the figure below.




Figure 1.2: The agent-environment interaction in a Markov Decision Process.

Based on environments, the rewards can also be randomly given according to the current state or the chosen action. In this sense, we generalize the transition probabilities into *dynamics*, which we denote by  $p(s', r | s, a)$ .

The agent and environment interact at each of a sequence of discrete time steps,  $t = 0, 1, 2, \dots$ . At each time step  $t$ , the agent receives its state  $S_t \in \mathcal{S}$  and takes action  $A_t \in \mathcal{A}$ . In part as a consequence of its action, the agent receives a reward  $R_{t+1} \in \mathcal{R}$ , and finds itself in a new state  $S_{t+1}$ . The MDP and agent together give rise to a stochastic process or trajectory as follows:

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$$

Based on the game, the episode can be either finite or infinite. If the state space of an MDP has a final state, then the episode terminates at the final state. In such cases, we usually write the trajectory as follows:

$$S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T, S_T$$

We also call the whole sequence an episode of the game/environment if it terminates at some time step. Notice that the states, actions, and rewards are written as random variables due to the dynamics and policy (we discuss about policy right in the next section). In this paper, we will assume that the state space and the action space is finite, that is,  $|\mathcal{S}| < \infty$  and  $|\mathcal{A}| < \infty$ .

## 2 Goal of the agent in an environment

The goal of an agent is to maximize the cumulative sum of rewards during the game. In this sense, given an MDP, the agent aims to **find a rule that chooses an action at each state**. We call such a rule the policy, which can be defined mathematically as a function  $\pi$  from the state space  $\mathcal{S}$  to the action space  $\mathcal{A}$ .

The policy  $\pi: \mathcal{S} \rightarrow \mathcal{A}$  can be either stochastic or deterministic. For example, in the grid world, a policy at the state (1,1) can be defined as follows:

$$\pi(a|s) = P(A_t = a|S_t = s) = \begin{cases} 0.6 & (a = \text{North}) \\ 0.4 & (a = \text{East}) \\ 0 & \text{otherwise} \end{cases} \quad (2.1)$$

The stochastic policy is indeed a conditional probability function. One should be able to discriminate two probabilities, the transition probability and the policy. Given a state  $S_t = s$ , the agent first chooses an action  $a$  according to the policy  $\pi(\cdot|s)$ .

The agent aims to find an optimal policy that maximizes the cumulative sum of rewards in an episode. This idea can be stated informally as the reward hypothesis.

**Remark 1** (Reward Hypothesis). *That all of what we mean by goals and purposes can be well thought of as the maximization of the expected value of the cumulative sum of rewards.*

We must stay tuned to the idea that the agent is not trying to maximize immediate reward but the cumulative sum of rewards in the long run. We define the cumulative sum of rewards formally as the return or the discounted return.

**Definition 2** (Return). *At each time step  $t$ , the return  $G_t$  is defined to be*

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots \quad (2.2)$$

*More generally, with the discount factor  $\gamma \in [0, 1]$  given in an MDP, the discounted return is defined to be*

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (2.3)$$

Note that the return is also a random variable. The discount rate determines the present value of future rewards: a reward received  $k$  time steps in the future is worth only  $\gamma^{k-1}$  times what it would be worth if it were received immediately. We usually use the discounted return for the following reasons.

1. It is mathematically convenient. If the rewards are bounded, then the discounted return converges.
2. The future is uncertain. That is, the agents do not know how many steps they should go to earn a specific future reward.
3. For agents, immediate rewards may earn more interest than delayed rewards.

Now we formally state the goal of the agents. The agent finds an *optimal policy* that maximizes the expected return  $\mathbb{E}_{\tau \sim \pi}[G_0(\tau)] = \int G_0(\tau) \pi(\tau) d\tau$ . Here,  $\tau$  is a random trajectory, and because of the Markov property, the probability  $\pi(\tau)$  of a terminating trajectory  $\tau$  is specified as follows:

$$\pi(\tau) = p(s_0) \prod_{t=0}^{T-1} \pi(a_t | s_t) p(s_{t+1} | s_t, a_t) \quad (2.4)$$

To summary, the agents seeks an optimal policy  $\pi^*$  which is equal to

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim \pi}[G_0(\tau)]. \quad (2.5)$$

### 3 Value Functions

Instead of maximizing the expected discounted return directly, we consider a method that assigns some scalar value for each state or each state-action pair. The value is determined by the expected return that an agent can receive in the environment starting in the state (and taking an action). The final goal of this method is also to obtain an optimal policy which maximizes the values. With keeping this concept in mind, we introduce the value functions.

**Definition 3** (Value Function). *Given an MDP and a policy  $\pi$  of an agent, the state-value function  $v_{\pi}(s)$  of a state  $s \in \mathcal{S}$  is defined to be the expected return when starting in  $s$  and following  $\pi$  thereafter. Formally,*

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]. \quad (3.1)$$

Here,  $\mathbb{E}_{\pi}[\cdot]$  denotes the expected value of a random variable given that the agent follows policy  $\pi$ , and  $t$  is any time step. The action-value function  $q_{\pi}(s, a)$  of taking action  $a \in \mathcal{A}$  in a state  $s \in \mathcal{S}$  is defined to be the expected return starting from  $s$ , taking the action  $a$ , and thereafter following  $\pi$ . Formally,

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a] \quad (3.2)$$

Note that, by the law of total probability, we have

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s] = \sum_{a \in \mathcal{A}} \pi(a | s) \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a] = \sum_{a \in \mathcal{A}} \pi(a | s) q_{\pi}(s, a). \quad (3.3)$$

A fundamental property of value functions used throughout reinforcement learning and dynamic programming is that they satisfy recursive relationships. Such recursive formulae are called the Bellman expectation equations.

**Theorem 4** (Bellman Expectation Equations). *The value functions can be decomposed into the immediate reward  $R_{t+1}$  and the discounted successor value function,  $\gamma v_\pi(S_{t+1})$  or  $\gamma q_\pi(S_{t+1}, A_{t+1})$ .*

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s] \quad (3.4)$$

$$q_\pi(s, a) = \mathbb{E}_\pi[R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a] \quad (3.5)$$

*Proof.* First we recall the definitions of value functions in Eq. 3.1 and Eq. 3.2.

$$v_\pi(s) = \mathbb{E}_\pi[G_t \mid S_t = s] \quad (3.1)$$

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a] \quad (3.2)$$

Starting from their definitions, we will derive the Bellman expectation equations. We first show the Bellman expectation equation for state-value function.

$$v_\pi(s) = \mathbb{E}_\pi[G_t \mid S_t = s] \quad (3.1)$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a) \quad (3.3)$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s) \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a] \quad (3.6)$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s) \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \quad (3.7)$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s', r} \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a, R_{t+1} = r, S_{t+1} = s'] \\ \times P(S_{t+1} = s', R_{t+1} = r \mid S_t = s, A_t = a) \quad (3.8)$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s']] \quad (3.9)$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')] \quad (3.10)$$

$$= \sum_{a \in \mathcal{A}} \sum_{s', r} p(a, s', r \mid s) [r + \gamma v_\pi(s')] \quad (3.11)$$

$$= \sum_{s', r} p(s', r \mid s) [r + \gamma v_\pi(s')] \quad (3.12)$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s] \quad (3.4)$$

Notice that Eq. 3.9 is obtained by Markov property. Next we show the Bellman expectation equation

for the action-value function.

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a] \quad (3.2)$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \quad (3.13)$$

$$= \sum_{r, s', a'} p(r, s', a' \mid s, a) \cdot \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a, R_{t+1} = r, S_{t+1} = s', A_{t+1} = a'] \quad (3.14)$$

$$= \sum_{r, s', a'} p(r, s', a' \mid s, a) \cdot \left[ r + \gamma \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s', A_{t+1} = a'] \right] \quad (3.15)$$

$$= \sum_{r, s', a'} p(r, s', a' \mid s, a) \cdot \left[ r + \gamma q_\pi(s', a') \right] \quad (3.16)$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a] \quad (3.5)$$

This completes the proof.  $\square$

We would like to note that, from Eq. 3.3 to Eq. 3.10, we have

$$q_\pi(s, a) = \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')]. \quad (3.17)$$

This equation tells you that the action-value function at  $(s, a)$  is the expectation of sum of the immediate reward  $r$  and the discounted state-value function at a successor state  $s'$ . Notice that the expectation is computed over the dynamics of the environment.

## 4 Optimal Policies and Optimal Value Functions

For finite MDPs, we can precisely define an optimal policy in the following way. Value functions define a partial ordering over policies. If we say that  $\pi \geq \pi'$  if and only if  $v_\pi(s) \geq v_{\pi'}(s)$  for all  $s \in \mathcal{S}$ , then the relation  $\geq$  on the set of policies becomes a partial order. We also say that  $\pi_*$  is an optimal policy if and only if  $\pi_* \geq \pi$  for any policy  $\pi$ . We will show that there is always at least one optimal policy in Theorem 8.

**Proposition 5.** *If  $\pi_*$  is an optimal policy, then for all state  $s \in \mathcal{S}$ ,*

$$v_{\pi_*}(s) = \max_{a \in \mathcal{A}} q_{\pi_*}(s, a). \quad (4.1)$$

Furthermore, we obtain that

$$v_{\pi_*}(s) = \max_{a \in \mathcal{A}} \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi_*}(s')]. \quad (4.2)$$

*Proof.* Observe that, for all  $s \in \mathcal{S}$ ,

$$v_{\pi_*}(s) = \sum_{a \in \mathcal{A}} \pi_*(a \mid s) q_{\pi_*}(s, a) \leq \max_{a \in \mathcal{A}} q_{\pi_*}(s, a).$$

For contradiction, assume that there is a state  $s'$  such that

$$v_{\pi_*}(s') < \max_{a \in \mathcal{A}} q_{\pi_*}(s', a) = q_{\pi_*}(s', a^*)$$

where  $a^* = \arg \max_{a \in \mathcal{A}} q_{\pi_*}(s', a)$ . Consider the policy  $\pi'$  such that

$$\pi'(a|s) = \begin{cases} \pi_*(a|s) & \text{if } s \neq s' \\ \begin{cases} 1 & \text{if } a = a^* \\ 0 & \text{otherwise} \end{cases} & \text{if } s = s' \end{cases}.$$

Then we have the following:

$$v_{\pi_*}(s') = \sum_{a \in \mathcal{A}} \pi_*(a|s') q_{\pi_*}(s', a) \quad (4.3)$$

$$< \sum_{a \in \mathcal{A}} \pi'(a|s') q_{\pi_*}(s', a) \quad (4.4)$$

$$= \sum_{a \in \mathcal{A}} \pi'(a|s') \sum_{s'', r} p(s'', r | s', a) [r + \gamma v_{\pi_*}(s'')] \quad (4.5)$$

$$= \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi_*}(S_{t+1}) | S_t = s'] \quad (4.6)$$

Next, we consider the term  $v_{\pi_*}(S_{t+1})$ . Since the state  $S_{t+1}$  may be the state  $s'$ , we obtain the following inequality:

$$v_{\pi_*}(S_{t+1}) = \sum_{a \in \mathcal{A}} \pi_*(a|S_{t+1}) q_{\pi_*}(S_{t+1}, a) \leq \sum_{a \in \mathcal{A}} \pi'(a|S_{t+1}) q_{\pi_*}(S_{t+1}, a). \quad (4.7)$$

Similar to the previous process through Eq. 4.4 to Eq. 4.6, we then obtain

$$v_{\pi_*}(S_{t+1}) \leq \sum_{a \in \mathcal{A}} \pi'(a|S_{t+1}) q_{\pi_*}(S_{t+1}, a) = \mathbb{E}_{\pi'} [R_{t+2} + \gamma v_{\pi_*}(S_{t+2}) | S_{t+1}]. \quad (4.8)$$

Thus, combining Eq. 4.6 and Eq. 4.8, we obtain the inequality:

$$\mathbb{E}_{\pi'} [R_{t+1} + \gamma v_{\pi_*}(S_{t+1}) | S_t = s'] \leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 v_{\pi_*}(S_{t+2}) | S_t = s'] \quad (4.9)$$

By continuing this process, we will obtain the inequality:

$$v_{\pi_*}(s') < \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s'] = v_{\pi'}(s') \quad (4.10)$$

This contradicts to the assumption that  $\pi_*$  is an optimal policy. Therefore, we conclude that, for all state  $s \in \mathcal{S}$ ,  $v_{\pi_*}(s) = \max_{a \in \mathcal{A}} q_{\pi_*}(s, a)$ . Furthermore, by Eq. 3.17, we also obtain

$$v_{\pi_*}(s) = \max_{a \in \mathcal{A}} \sum_{s', r} p(s', r | s, a) [r + \gamma v_{\pi_*}(s')].$$

This completes the proof.  $\square$

Proposition 5 is also called the Bellman operator equations. This will be again introduced after we define what optimal value functions are.

**Definition 6** (Optimal Value Function). *The optimal state-value function  $v_*$  and the optimal action-value function  $q_*$  is defined as*

$$v_*(s) = \max_{\pi} v_{\pi}(s), \quad \forall s \in \mathcal{S}, \quad (4.11)$$

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a), \quad \forall s \in \mathcal{S}, \forall a \in \mathcal{A}. \quad (4.12)$$

**Proposition 7.** *For any two optimal policies  $\pi_1$  and  $\pi_2$ , we have  $v_{\pi_1}(s) = v_{\pi_2}(s)$  for all  $s \in \mathcal{S}$ .*

*Proof.* Since  $\pi_1$  is optimal, we have  $v_{\pi_1}(s) \geq v_{\pi_2}(s)$  for all  $s \in \mathcal{S}$ . Likewise, since  $\pi_2$  is optimal, we have  $v_{\pi_2}(s) \geq v_{\pi_1}(s)$  for all  $s \in \mathcal{S}$ . This shows that  $v_{\pi_1}(s) = v_{\pi_2}(s)$  for all  $s \in \mathcal{S}$ .  $\square$

By Proposition 7, we can see that all the optimal policies share the same state-value function. According to the definition of optimal policies, for any optimal policy  $\pi_*$ , we can see that

$$v_{\pi_*}(s) = \max_{\pi} v_{\pi}(s) = v_*(s)$$

for all state  $s \in \mathcal{S}$ . However, since the maximization in the definition of optimal value functions is done independently of each state, the optimal policies  $\arg \max_{\pi} v_{\pi}(s)$  and  $\arg \max_{\pi} v_{\pi}(s')$  may not be the same if  $s \neq s'$ . In other words, an optimal policy that maximizes the state value function at a state  $s$  may not be the same policy that maximizes the state value function at another state  $s'$ . With this caution in mind, consider the following theorem.

**Theorem 8** (Optimal Policy and Optimal Value Functions). *Any finite MDP satisfies the following.*

1. *There exists an optimal policy  $\pi_* \geq \pi$  for all  $\pi$ .*
2. *All optimal policies achieve the optimal state-value function  $v_{\pi_*}(s) = v_*(s)$  for all  $s \in \mathcal{S}$ .*
3. *All optimal policies achieve the optimal action-value function  $q_{\pi_*}(s, a) = q_*(s, a)$  for all  $(s, a) \in \mathcal{S} \times \mathcal{A}$ .*

*Proof.* Assuming the existence of optimal policies, the second and the third statements are immediate by Proposition 7. Let a state  $s \in \mathcal{S}$  be given. Choose an optimal policy  $\pi_s$  such that

$$\pi_s = \arg \max_{\pi} v_{\pi}(s).$$

For any optimal policy  $\pi_*$ , we have  $v_{\pi_*}(s) = v_{\pi_s}(s)$  by Proposition 7. Hence,

$$v_*(s) = \max_{\pi} v_{\pi}(s) = v_{\pi_s}(s) = v_{\pi_*}(s).$$



This shows that the second statement holds. Now we prove the third statement. For an optimal  $\pi_*$ ,

$$\begin{aligned}
q_{\pi_*}(s, a) &= \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi_*}(s')] && \text{(by Eq. 3.17)} \\
&= \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')] \\
&= \sum_{s', r} p(s', r \mid s, a) [r + \gamma \max_{\pi} v_{\pi}(s')] \\
&= \max_{\pi} \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')] && (\because |\mathcal{S}| < \infty, |\mathcal{A}| < \infty) \\
&= \max_{\pi} q_{\pi}(s, a) && \text{(by Eq. 3.17)} \\
&= q_*(s, a)
\end{aligned}$$

This shows the third statement. Now it remains to prove the existence of an optimal policy. We first define the Bellman optimality operator.

**Definition 9** (Bellman Optimality Operator). *The Bellman optimality operator  $\mathcal{T}: \mathbb{R}^{|\mathcal{S}|} \rightarrow \mathbb{R}^{|\mathcal{S}|}$  is an operator that satisfies the following: for any vector  $V = \{v(s) : s \in \mathcal{S}\} \in \mathbb{R}^{|\mathcal{S}|}$ ,*

$$(\mathcal{T}V)(s) = \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma V(s')] \quad (4.13)$$

Notice that, for an optimal policy  $\pi_*$ , we have  $\mathcal{T}v_{\pi_*} = v_{\pi_*}$ .

**Definition 10** (Contraction). *Let  $X$  be a metric space, with metric  $d$ . If  $\varphi$  maps  $X$  into  $X$  and if there is a number  $c < 1$  such that*

$$d(\varphi(x), \varphi(y)) \leq c d(x, y)$$

*for all  $x, y \in X$ , then  $\varphi$  is said to be a contraction of  $X$  into  $X$ .*

We will show that the Bellman optimality operator  $\mathcal{T}$  is a contraction mapping in infinity norm,

$$\|V\|_{\infty} = \max_s |v(s)|$$

and we use the Banach fixed-point theorem to show the existence of optimal policy.

**Theorem 11** (Banach fixed-point theorem). *Let  $(X, d)$  be a non-empty complete metric space with a contraction mapping  $\mathcal{T}: X \rightarrow X$ . Then  $\mathcal{T}$  admits a unique fixed-point  $x^*$  in  $X$ , i.e.,  $\mathcal{T}(x^*) = x^*$ . Furthermore,  $x^*$  can be found as follows: start with an arbitrary element  $x_0$  in  $M$  and define a sequence  $\{x_n\}$  by  $x_n = \mathcal{T}(x_{n-1})$  for  $n \geq 1$ . Then  $x_n \rightarrow x_*$  exponentially rapidly.*

Before we show that the Bellman optimality operator  $\mathcal{T}$  is a contraction, we state a lemma.

**Lemma 12.** *For any two real-valued functions  $f, g: \mathcal{A} \rightarrow \mathbb{R}$ ,*

$$\left| \max_a f(a) - \max_a g(a) \right| \leq \max_a |f(a) - g(a)| \quad (4.14)$$

*Proof.* Without loss of generality, assume that  $\max_a f(a) \geq \max_a g(a)$ . Let  $a^* = \arg \max_a f(a)$ . Then,

$$\left| \max_a f(a) - \max_a g(a) \right| = \max_a f(a) - \max_a g(a) = f(a^*) - \max_a g(a) \leq f(a^*) - g(a^*) \leq \max_a |f(a) - g(a)|$$

This proves the lemma.  $\square$

Now we show that the Bellman optimality operator  $\mathcal{T}$  is a contraction. For any  $V, V' \in \mathbb{R}^{|S|}$ ,

$$\|\mathcal{T}V - \mathcal{T}V'\|_\infty = \max_s |\mathcal{T}V(s) - \mathcal{T}V'(s)| \quad (4.15)$$

$$= \max_s \left| \max_a \sum_{s',r} p(s', r|s, a)[r + \gamma V(s')] - \max_a \sum_{s',r} p(s', r|s, a)[r + \gamma V'(s')] \right| \quad (4.16)$$

$$\leq \max_s \max_a \left| \sum_{s',r} p(s', r|s, a)[r + \gamma V(s')] - \sum_{s',r} p(s', r|s, a)[r + \gamma V'(s')] \right| \quad (4.17)$$

$$\leq \max_s \max_a \left| \sum_{s'} p(s'|s, a) \gamma (V(s') - V'(s')) \right| \quad (4.18)$$

$$\leq \gamma \max_{s'} |V(s') - V'(s')| \quad (4.19)$$

$$= \gamma \|V - V'\|_\infty \quad (4.20)$$

Banach-fixed theorem implies that, there is a unique fixed-point  $V^*$ , that is,  $\mathcal{T}V^* = V^*$ . It still remains to show that  $V^*$  is a value function for some policy  $\pi$ . If we show that  $V^* = v_\pi$  for some  $\pi$ , then  $\square$

To summarize, the optimal value functions  $v_*$  and  $q_*$  are value functions corresponding to any optimal policy  $\pi_*$ . Therefore, the uniqueness of the optimal value functions implies that the Bellman expectation equation for an optimal policy can be restated in a special form without reference to any specific policy. We call the equation *Bellman optimality equation*.

**Theorem 13** (Bellman Optimality Equation).

$$v_*(s) = \max_a \sum_{s',r} p(s', r|s, a)[r + \gamma v_*(s')] \quad \left( = \max_a q_*(s, a) \right) \quad (4.21)$$

$$q_*(s, a) = \sum_{s',r} p(s', r|s, a)[r + \gamma \max_{a'} q_*(s', a')] \quad \left( = \sum_{s',r} p(s', r|s, a)[r + \gamma v_*(s')] \right) \quad (4.22)$$

## References

- [1] Andrew Barto and Richard S. Sutton, *Reinforcement Learning: An Introduction* (2nd ed.), The MIT Press, 2018.
- [2] Alireza Modirshanechi, *Why does the optimal policy exist?* Medium, towards data science blog, <https://towardsdatascience.com/why-does-the-optimal-policy-exist-29f30fd51f8c>